

Assembly Language Specification

2022 Spring, SWPP

1. Architecture Overview

- An architecture consists of a single-core CPU and 64-bit memory space.

(1) Registers

- There are 33 64-bit general registers. They are named r1, r2, ..., r32, and sp.
- r1, r2, ..., r32 are initialized to 0, and sp is initialized to 102400.
- A register can be assigned multiple times (it isn't SSA).

(2) Memory

Loads and stores.

- The memory is accessed via load/store/aload/await instructions with 64-bit pointers.
- The exact formula for its calculation is described later.
- An asynchronous memory load can be performed using aload (asynchronous load) and await (waiting until asynchronous to be resolved).

Stack.

- The stack area starts from address 102400, grows downward (-), and is initialized as 0 at the beginning of the program execution.
- You can use sp to store the address of the current stack frame, but it is not necessary to do so.

Heap.

- The heap area starts from address 204800 and grows upward (+).
- Heap allocation (malloc) initializes the area as zero.
- Accessing an unallocated heap raises an error.
- Accessing the area between [102400, 204800) raises an error.

Global Variables.

- Syntactically, there is no difference between global variables and heap-allocated blocks.
- The project skeleton lowers a global variable to a heap allocation (malloc call) at the beginning of the main(). So, they are placed at the beginning of the heap area.

(3) Function calls

- Function arguments can be accessed via read-only registers arg1.. arg16.

Calling convention.

- When a call instruction is executed,
 - r1 ~ r32, sp registers are automatically saved in an invisible space (you don't need to manually spill them).
 - Values of the arguments are automatically assigned to the registers arg1 ~ arg16.
 - The values of r1 ~ r32 are unchanged (not initialized to 0).
- After the call returns, r1 ~ r32, sp registers are automatically restored.

(4) Cost

- The execution cost of a program can be calculated as 'program-wide instruction execution cost + maximum heap memory usage (in bytes)'.
- The code size is irrelevant to the total cost.

Memory usage cost.

- The memory usage cost is 16 times the maximum heap-allocated byte size at any moment.
- For example, the memory usage cost of

```
r1 = malloc 8
free r1
r2 = malloc 8
free r2
```

is $16 * 8 = 128$, because the maximum memory usage is 8 bytes.

Compile time.

- Compile time should be less than 1 minute.

2. Input Program

Structure.

- The source program consists of a single IR file; There is no linking.
- The IR file consists of one or more functions, including the main function.
- A source program only uses i1, i8, i16, i32, i64, array types, and pointer types.

Function.

- A function can have at most 16 arguments.
- There is no function attribute (e.g. read-only).
- `main()` is never called recursively.

Standard I/O.

- A source program takes input through `read()` calls. `read()` reads an integer and returns it as a 64 value.
- The output of the program is done via `write(i64)` calls. It writes the output as an unsigned integer in a new line.
- `read()` / `write(i64)` calls are connected to the standard input/output.

Misc.

- The test programs will never raise out-of-memory or stack overflow with the given inputs if compiled with the project skeleton.

3. Function & Basic Block

(1) Function

Syntax:

```
start <funcname> <Narg>:  
    ... (basic blocks)  
end <funcname>
```

- A function contains one or more basic blocks.
- <funcname> is a non-empty string consisting of alphabets(a-zA-Z), digits(0-9), underscore(_), hyphen(-), or dot(.).
- <Narg> describes the number of arguments.
- A function's return type is always i64.
- There is no variadic function.
- There is no nested function.

(2) BasicBlock

Syntax:

```
<bbname>:  
    ... (instructions)
```

- A basic block consists of one or more instructions.
- A basic block must end with a terminator instruction (see below for more details)
- <bbname> is a non-empty string, starting with a dot(.) and consists of alphabets(a-zA-Z) + digits(0-9) + underscore(_) + hyphen(-) + dot(.).

(3) Comment

Syntax:

```
; <comment>
```

- A comment starts with a semicolon(;).
- Only space characters are allowed before the semicolon in the line.

4. Instructions

Syntax:

```
op_name <val1> .. <valN>
<reg> = op_name <val1> .. <valN>
```

- `<reg>` is the name of a register to assign the result.
- `<val>` is one of integer constant, `bbname`, or a register. `<valk>` is the k -th operand of the instruction.
- Argument registers (e.g. `arg1`) cannot be placed at the LHS.

(1) Terminator instructions

Kind	Syntax	Cost
Return Value - ret is equivalent to ret 0.	ret ret <val>	1
Unconditional Branch	br <bbname>	1
Conditional Branch	br <condition> <>true_bb> <>false_bb>	6 for true_bb 1 for false_bb
Switch Instruction - <val1>, ... should be constant integers.	switch <cond_val> <val1> <bb1> .. <default_bb>	1.2

- Terminator instructions should come at the end of a basic block only.
- `<bbname>` stands for a basic block name to jump to.
- Branches/switch cannot jump to a block in another function.
- `ret` does not reset the temperature of the previously used stack area.

(2) Memory allocation/deallocation

Kind	Syntax	Cost
Heap Allocation	<reg> = malloc <val>	16
Deallocation	free <reg>	16

malloc.

- malloc allocates a space of the given size to the heap. The space is initialized to zero and has zero temperature.
- The size of malloc should be non-zero & a multiple of 8.
- malloc finds an empty consecutive space with the smallest address in the heap area & allocates it.
- The returned address by malloc is a multiple of 8.

free.

- free deallocates a space associated with the given pointer.
- The pointer passed to free should point to the beginning of allocated heap space.

(3) Memory access

Kind	Syntax	Base Cost
Load	<code><reg> = load <size> <ptr></code> <code><size> := 1 2 4 8</code>	Stack area: 6 Heap area: 12
Store	<code>store <size> <val> <ptr></code> <code><size> := 1 2 4 8</code>	Stack area: 6 Heap area: 12
Async Load	<code><reg> = aload <size> <ptr></code> <code><size> := 1 2 4 8</code>	Stack area: 1 Heap area: 1
Await waits for cost n until the async load is resolved.	<code>await <reg></code>	$1 + (0 \sim n)$ Stack area: $n = 10$ Heap area: $n = 16$

load.

- The load instruction reads the data at [`<ptr>` , `<ptr>+<size>`), zero-extends it to 64 bits, and returns it.
- `<ptr>` should be multiple of `<size>`.
- The memory is *little-endian*. The least significant byte of the value read by load is from `<ptr>`, and the most significant byte is from `<ptr>+<size>-1`.

store.

- The store instruction truncates the value `<val>` to an `<size>*8-bit` integer and writes it at [`<ptr>`, `<ptr>+<size>`).
- `<ptr>` should be multiple of `<size>`.

asynchronous load and await.

- The `aload` instruction behaves exactly the same as an ordinary load, except that the returning register is not usable immediately after executing the `aload` instruction. Instead, one should

wait until the async load is resolved using await <reg> instruction. e.g.,

```
r2 = aload 8 r1
...      ; any use (read/write) of r2 is an error
await r2 ; waits until r2 is ready
...      ; can use r2
```

- After executing an aload instruction, it takes the cost $n = 10$ for stack area and $n = 16$ for heap area for the returning register to be ready. e.g.,

```
r2 = aload 8 r1 ; suppose r1 contains an address to heap area
...            ; instructions here take cost  $m$ 
await r2       ; waits for  $16 - m$ 
```

Total cost = 1 (for aload) + $m + 1$ (for await) + $\max(0, 16 - m)$ (for waiting)

- Accesses to addresses overlap an unresolved async load is allowed. e.g.,

```
store 8 3 r1
r2 = aload 8 r1
store 8 42 r1
r3 = load 8 r1 ; loads 42
await r2       ; r2 = 3 when resolved
```

- An await that is not paired with a previous async load is allowed, and the cost of it is always 1

(4) Other instructions

Kind	Name	Cost
Integer Multiplication/Division	<reg> = udiv <val1> <val2> <bw> <reg> = sdiv <val1> <val2> <bw> <reg> = urem <val1> <val2> <bw> <reg> = srem <val1> <val2> <bw> <reg> = mul <val1> <val2> <bw> <bw> := 1 8 16 32 64	1
Integer Shift/Logical Operations - shl: shift-left - lshr: logical shift-right - ashr: arithmetic shift-right	<reg> = shl <val1> <val2> <bw> <reg> = lshr <val1> <val2> <bw> <reg> = ashr <val1> <val2> <bw> <reg> = and <val1> <val2> <bw> <reg> = or <val1> <val2> <bw> <reg> = xor <val1> <val2> <bw> <bw> := 1 8 16 32 64	4
Integer Add/Sub	<reg> = add <val1> <val2> <bw> <reg> = sub <val1> <val2> <bw> <bw> := 1 8 16 32 64	5
Integer Sum	<reg> = sum <val1> ... <val8> <bw> <bw> := 1 8 16 32 64	5.2

Integer increment <reg> = <reg> + 1	incr <reg> <bw>	1
Integer decrement <reg> = <reg> - 1	decr <reg> <bw>	1
Comparison - <cond> is equivalent to the cond of LLVM IR's icmp	<reg> = icmp <cond> <val1> <val2> <bw> <bw> := 1 8 16 32 64	1
Ternary operation	<reg> = select <val_cond> <val_true> <val_false>	1.2
Function call	call <fname> <val1> .. <valN> <reg> = call <fname> <val1> .. <valN>	2 + arg #
Assertion An assertion fail is an error Used for testing	assert_eq <val1> <val2>	0

- For integer arithmetic and comparison operations, <bw> is the size of bitwidth of inputs that the operation assumes. For example, `ashr 511 2 8` takes the lowest 8-bits from inputs (which is 255 = -1), performs arithmetic right shift, and zero-extends it to 64 bits. So, its result is 255.